

Limited Direct Execution

Idea: To run as fast as possible, user code runs directly on the CPU.

But...

Some instructions are too

- powerful
- alter machine state beyond one process

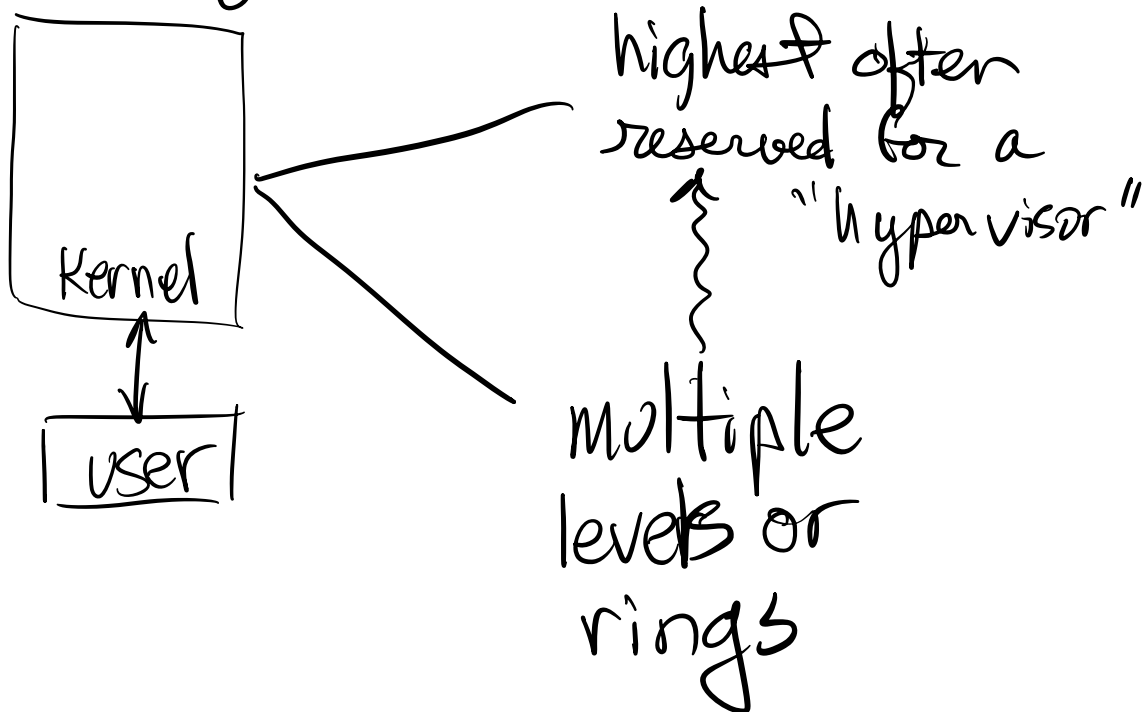
Therefore...

Some instructions, if attempted by user process, kill the process.

So how do these special instructions get executed?

AND... how to enter/exit the kernel state - for example - system calls.

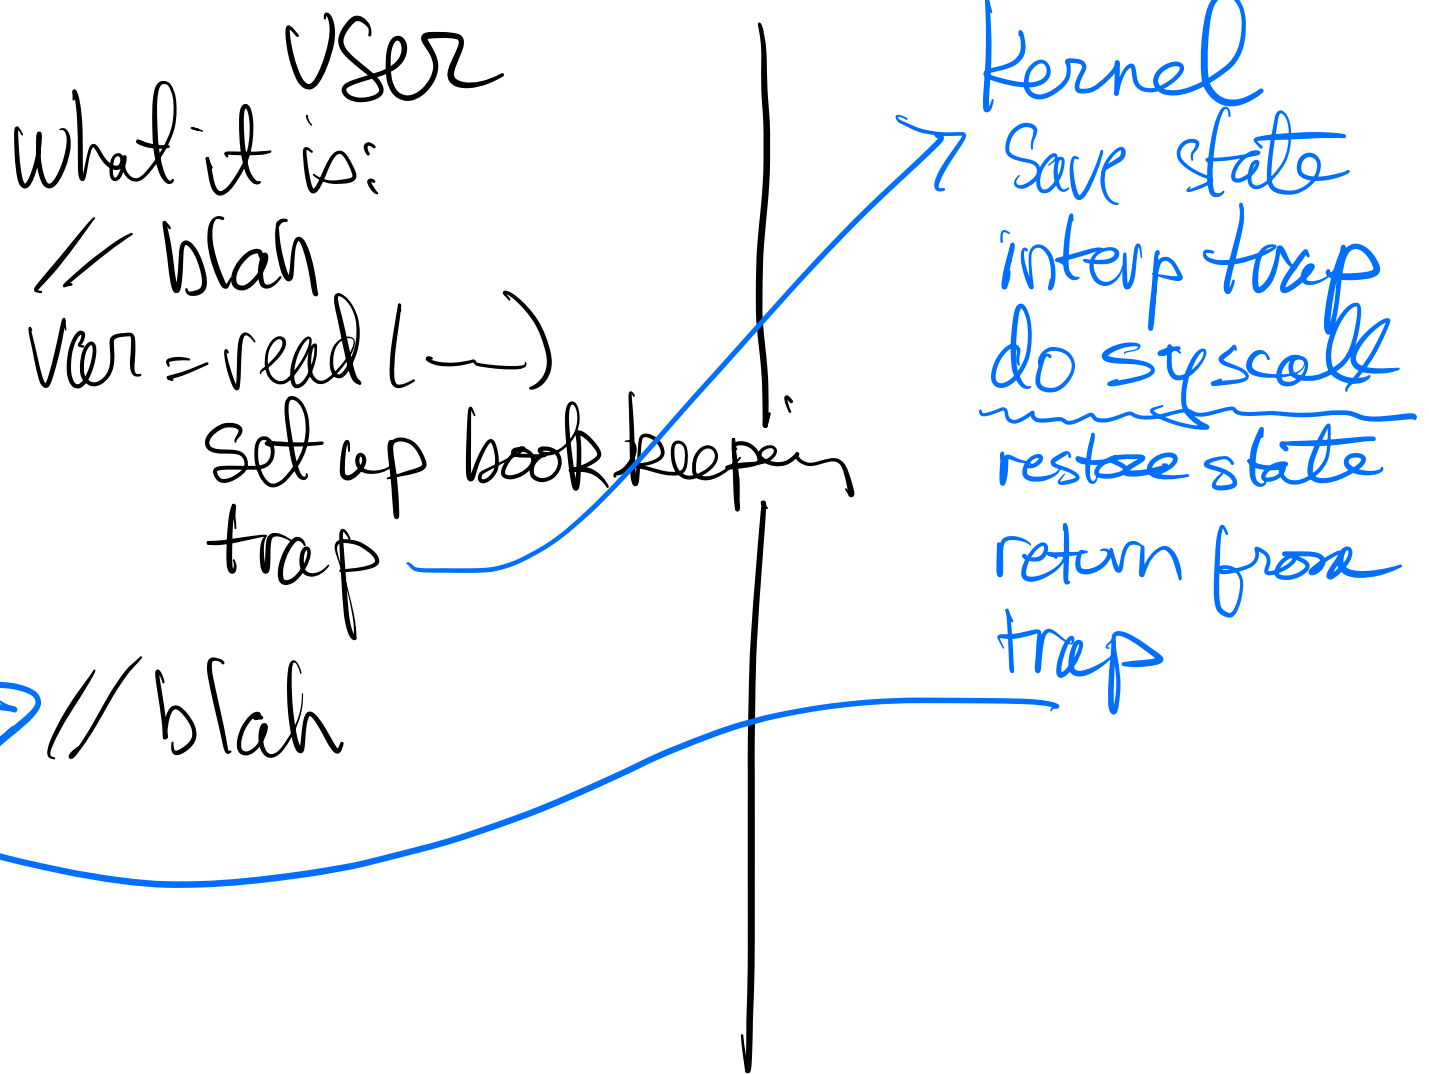
User processes cannot access kernel memory. Solution: process states



When a user process needs the services of the kernel, it TRAPS to cause a state transition.

What it looks like:

```
// blah blah blah
    var = read(_____);
// blah blah blah
```



The open() system call

open(...) → put parameters in right place

put "I am a system call" in right place

put "I am Open" in right place

trap.

ARM64

syscall number \rightarrow X8

parameters \rightarrow X0 \rightarrow X7

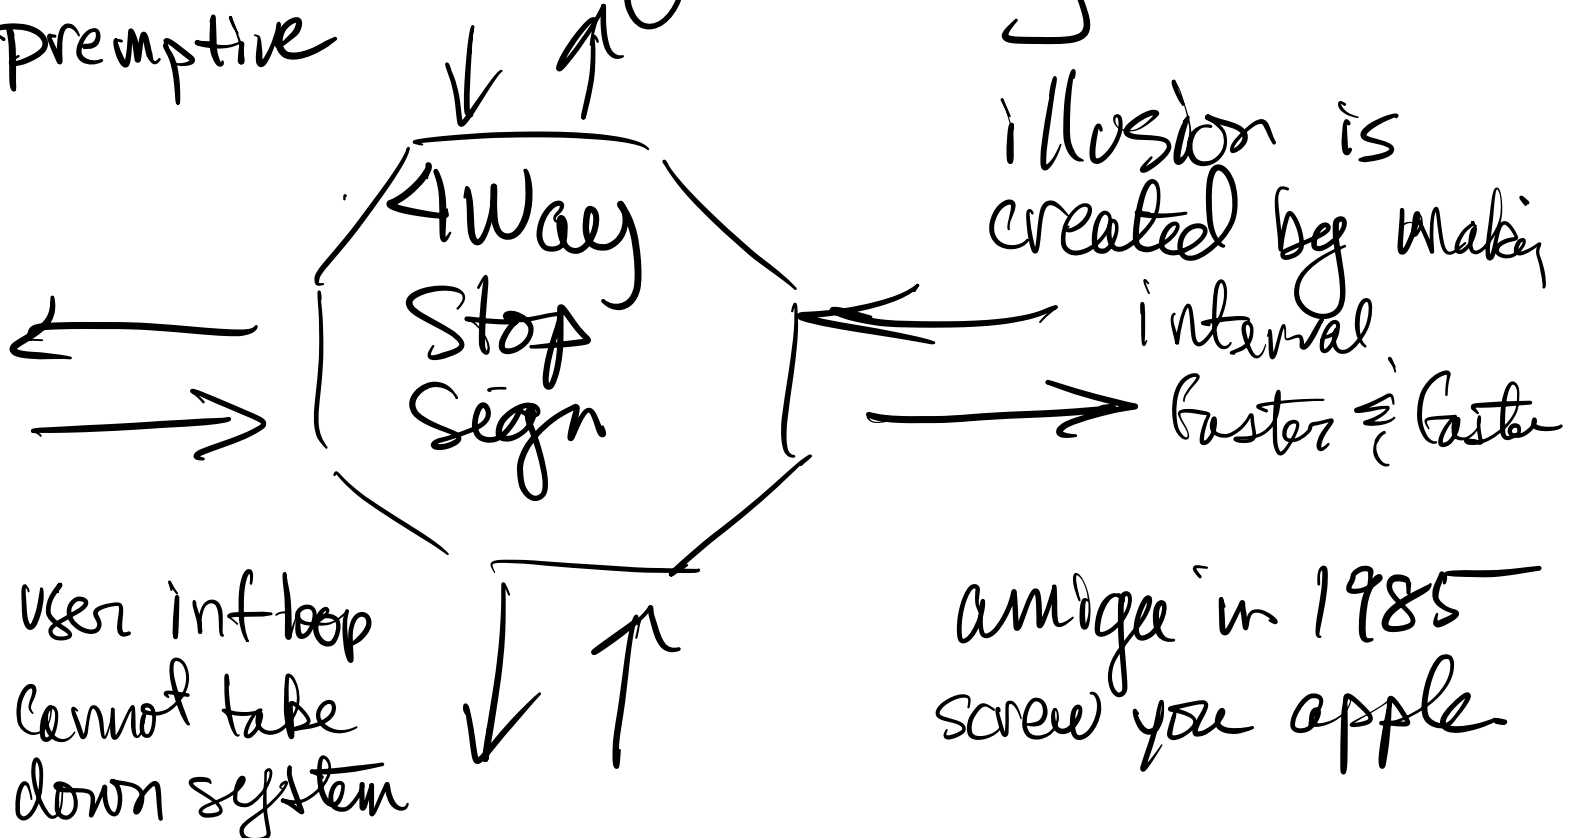
trap \rightarrow SVC \emptyset

return value in X \emptyset

More details on RiscV look at
trampoline.S uservec

Methods of multitasking
— illusion of owning the CPU

Preemptive



cooperative

yield() or sleep() etc
make a system call
in general.

Mac OS 1 \rightarrow 9 till 2001

1: b 1b or in other
words

label: branch to label.

Any kind of infinite loop in user code
causes tight lock up!!!

Context switch $\approx 1 \mu s$
equivalent to about 2000 instructions