

MLFQ

multi level feedback queue.

In situations where you can adaptively learn from history, do so.

Goal of MLFQ

Keep interactive processes interactive

Original Unix scheduler - i.e. the one in the original V6

The more you run, the lower your priority

The lower your priority, the less you run.

The less you run, the higher your priority

Compute
band

IO
band

Interactive processes (e.g. V.I)
spends most of its time waiting for
input ... not running ... so priority
stays high!

MLFQ

assumptions:

- * Multiple levels of priority.
- * Each has a queue of runnables
- * Running tasks are not on any queue
- * Waiting / blocked tasks are not
* on any queue.

Rule 1 if $Pri(A) > Pri(B)$
run A ... ie A is on
queue above B.

Rule 2 RR^* for tasks on same queue
* or other algorithm.

Rule 3 Jobs arrive on top queue

Rule 4a When job uses up quantum it is demoted.

4b If job gives up CPU, it returns to same queue (eg I/O)





Figure 8.3: Along Came An Interactive Job

See how short job finished quickly?
even w/o knowing ahead of time?



Doing I/O kept job in gray at high priority because it ran only briefly.
 Could be an editor for example

Notice so far jobs only sink in priority.
 What if compute bound job changes to interactive? tough

Rule 5

After a while (called EPOCH)
put every job back on
highest priority queue and
start over.

Revise Rule 4 - no more oppty
to cheat. Time on cpu measured
cumulatively - can't cheat by doing
an I/O \triangleleft

MLFQ is basis scheduler for
BSD derivatives

Solaris

Windows NT \rightarrow Win 11